


Deep Learning and Practice

Lab2: Temporal Difference Learning

H092093 林佳榮

- A plot shows episode scores of at least 100,000 training episodes



▲ Episode scores of 100,000 training episodes

To plot the episode scores chart, I took the mean value of every 1000 episodes as a unit. Therefore, there are totally 100 points on the chart up above which indicates the average value of every 1k episodes from episode 1 to episode 100,000.

- Describe the implementation and the usage of n-tuple network

According to TA's introduction on n-tuples, the reason we use n-tuple networks is because it is impossible to record all the state's and the value respectively.

The board of 2048 is 4×4 , which means that there would be 2^{64} possibilities, which requires too much memory. So if we use 4 kinds of 6-tuple network, it can be $4 \times 15^6 \times 4$, which can greatly reduce the memory usage.

In the given sample code, a weight table is defined in the feature class and is initialized to later on will be used to get the state value. Notably, 8 isomorphisms of a pattern share the same weight table.

- **Explain the mechanism of TD(0)**


The mechanism of TD(0) is to update the state value by reward and value function *every time step, meaning don't have to wait until the episode is over, it can directly use the next state value to do an update*. Compared to Monte-Carlo Backup(another model-free learning method), which has to wait until an episode is over to get the return value to update state value:

$$\text{MC update: } V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

TD backup update value every single step:

$$\text{TD backup: } V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + V(S_{t+1}) - V(S_t))$$

- **Explain the TD-backup diagram of V(after-state)**



While we are doing TD-learning, after an episode is over, we have to do the updating process: from terminal to initial state, we update every estimated state value by another estimated value on the next time step, which is a kind of bootstrapping method. There are total 2 different updating methods: (before-)state, after-state TD-backup, classified by which state value is targeted to be updated:

In the after-state version, the value of **after-state** is updated with the value of **after-state** at the next time step.

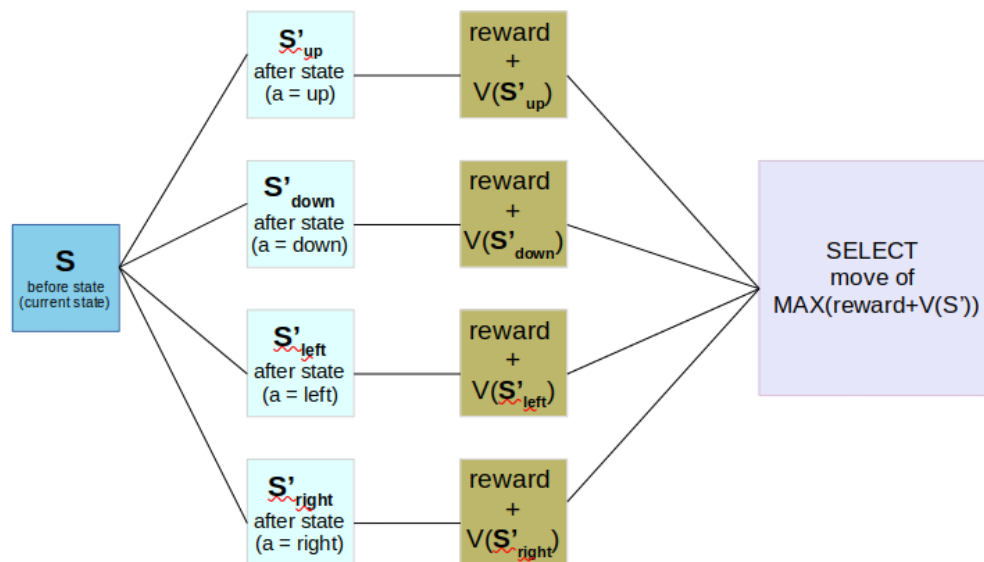
This is how it is backedup:

$$td_error: (reward_{next} + V(s'_{next})) - V(s')$$

$$V(s') = V(s') + \alpha * td_error$$

where alpha is the learning rate (step-size), $V(s')$ is the current estimated after state's value and $V(s'_{next})$ is the estimated next after-state's value. So in the above diagram the value of the blue box S' will be updated based on the value of the blue of S'_{next} on the next time step.


- Explain the action selection of $V(\text{after-state})$ in a diagram



▲ Diagram 1: after-state version action selection process

Diagram 1 shows briefly the workflow of the after-state version in selecting the best move. Here is the explanation:

Let's start with the following example:



1. We are at the current (before)-state S . In order to choose the best next move, we have to consider the following four cases: moving up/down/left/right.
2. Then, we take the action, moving up/down/left/right which makes the current (before)-state transfer to the current after-state s' , that's to say, we'll have four corresponding after states: $S'_{up} / S'_{down} / S'_{left} / S'_{right}$
3. Thus, we calculate the corresponding value of reward + value of current_after_state $V(s')$, so we'll have:

$$\text{value}_1 = \text{reward} + V(S'_{up})$$

$$\text{value}_2 = \text{reward} + V(S'_{down})$$


$$\text{value}_3 = \text{reward} + V(S'_{left})$$

$$\text{value}_4 = \text{reward} + V(S'_{right})$$

4. Finally, we'll choose the move which forms the maximum value among the four.

$$\text{best move} = \text{argmax}_{\text{move} = \{\text{up, down, left, right}\}}(\text{value})$$

- Explain the TD-backup diagram of $V(\text{state})$



▲ Diagram 2: State version TD backup diagram

Similar to TD-backup diagram of $V(\text{after-state})$, there are only slight differences when it comes to state-version td-backup. The significant difference is the object that is updated. In the state version, the value of **before-state** is updated with the value of **before-state** at the next time step.

This is how it is backuped:

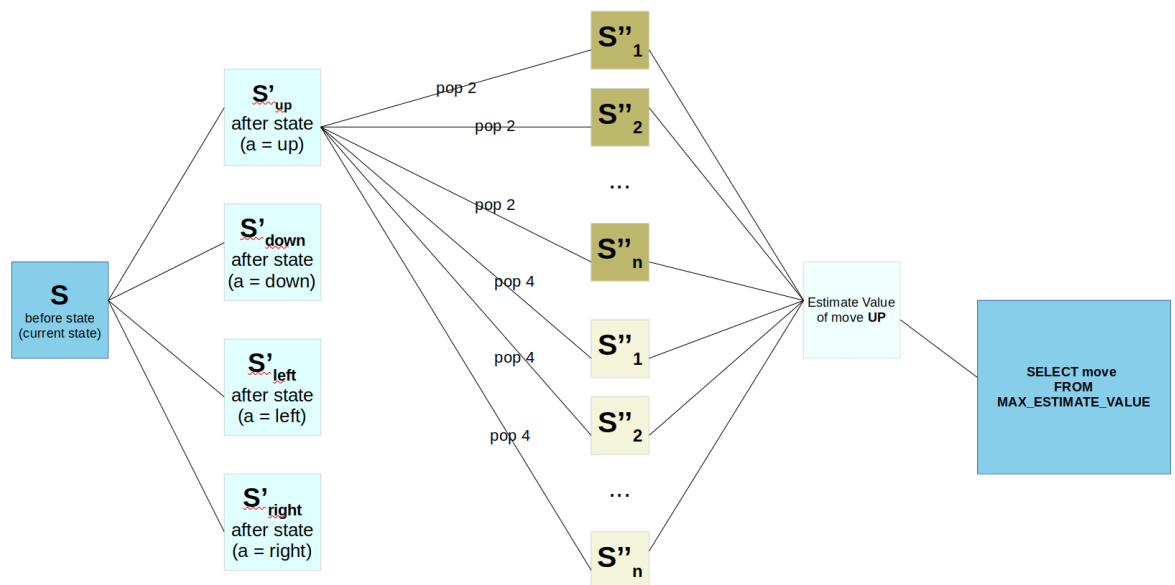
$$\text{td_error: } (\text{reward} + V(s'')) - V(s)$$

$$V(s) = V(s) + \alpha * td_error$$

where α is the learning rate (step-size), $V(s)$ is the current estimated state's value and $V(s')$ is the estimated next state's value. Therefore, we can tell that the difference between TD-backup diagrams of $V(\text{state})$ and $V(\text{after-state})$ is how it is backedup:

State version's td error to calculate the difference between *reward + next_before_state_estimated_value* and *current_before_state_estimated value*, while after-state version will calculate the difference between *next_state_reward + next_after_state_estimated_value* and *current_after_state_value*.

- Explain the action selection of $V(\text{state})$ in a diagram




▲ Diagram 3: state version action selection process


The diagram up above shows briefly how state version in selecting the best move works and here is the explanation:

Let's take the following board for example:

1. Board on the left hand side is where we are now. It is the before state S , which is our current state. In order to determine which move we are going to take, we have to test all move possibilities: moving up/down/left/right, and calculate the corresponding estimation value to decide the best move. For example, we test moving up first.




▲ Current state : **S** Before-state
after-state



▲ Moving Up: **S'**

2. After we take action of moving up, in after state S' , there will be 2×14 possible after-state boards, because there may either pop 2 or 4 in any blanks in the S' board, which leads to the next state S'' . Therefore, we have to enumerate every kind of possibility of next state S'' , like in **Diagram 3** the brown cubes indicate, here list the possibilities:



▲ Possibilities of all next-state S'' : after-state S' + (pop2 or pop4)

After all the next-states are listed, we then calculate the corresponding estimation value according to this formula:

$$\sum_{s'' \in S''} P(s, a, s'') V(s'')$$

so we'll have

pop 2: $0.9 * estimate(next_state_board)$

pop 4: $0.1 * estimate(next_state_board)$ ¹

Therefore, the estimation value of making this move will be:

estimation_value_of_the_move = reward + (pop2 + pop4) / numbers_of_blank

By calculating the estimation value of making each move (up/down/left/right), we can then select the best move by choosing the move which leads to the maximum estimation value:

best move = $\underset{action = \{up, down, left, right\}}{argmax} (estimation_value)$

- **Describe your implementation in detail**

To train to play the 2048 game this time, there are a total of 5 TODOs within the given sample code. To describe my implementation in detail, I'll explain what I've done separately and sequentially with algorithm-like descriptions.

I implemented TODO1, 2, 3 based on the sample github, so in the following description I'll focus more on TODO 4 and 5.

TODO 1: estimate

To estimate the given board's value, we have to look up the weight table and calculate the given board's patterns and its isomorphism value.

```
for (i = 0 to last_isomorphism):  
    //look up the weight table  
    value += indexof(isomorphic[i], b)'s value
```

TODO 2: update

In this part, we have to update 8 isomorphic feature values of specific patterns and return the sum of the updated value to make it the state's new value.

TODO 3: indexof

¹ According to 2048's rule, the probability of popping 2 is 0.9 and popping 4 is 0.1.

Since I use the default setting of 6-tuple network with each tile represented by 4 bit, so getting index will be:

$$\text{index} \mid= \text{b.at}(\text{pattern}[\text{i}]) \ll (4 * \text{i})$$

TODO 4: select best move

Here is the pseudo code how I implement the select_best_move part:

```
for ( move = up, down, left, right):  
  
    all_possible_board = move.after_state()  
  
    for (i=0 to 15):  
        if all_possible_board.at(i) == 0:  
            blank+=1  
            all_possible_board.set(i, 1)  
            pop_two += 0.9 * estimate(all_possible_board)  
            all_possible_board.set(i, 2)  
            pop_four += 0.1 * estimate(all_possible_board)  
            all_possible_board.set(i, 0)  
  
    sum = (pop_two + pop_four)/blank  
    move->set_value(reward + sum)  
  
    if move->value() > best->value():  
        best = move
```

The goal of the step is to pick the move with the best estimation value by considering the four cases: moving up/down/left/right by iterating through the for loop.

After picking a move:

First, record the after-state-board to all_possible_board:

$$\text{all_possible_board} = \text{move.after_state}()$$

Second: **for(i=0 to 15)** enumerate every kind of next state possibilities based on current move by iterating through the board and popping up 2 or 4 at any blank.

About how to find a blank: if the board's tile is 0, it means it's a blank at that place.

If it is a blank, then pop either 2 or 4 at that place:

Case 1: pop 2 on all_possible_board, then calculate the estimation value by $0.9 * \text{board_after_pop_2's_value}$:

```
all_possible_board.set(i, 1)  
pop_two += 0.9 * estimate(all_possible_board)
```

Case 2: pop 4 on all_possible_board, then calculate the estimation value by $0.1 * \text{board_after_pop_4's_value}$

```
all_possible_board.set(i, 2)  
pop_four += 0.1 * estimate(all_possible_board)
```

After calculation, the blank with 2 or 4 popped must be set back to 0.

```
all_possible_board.set(i, 0)
```

Value is set based on this formula: $\text{Value} = r + \text{Sigma}(P(s,a,s'')V(S''))$, where $\text{Sigma}(P(s,a,s'')V(s''))$ is the sum here

```
sum = (pop_two + pop_four)/blank  
move->set_value(reward + sum)
```

if current move brings bigger value, then it is better:

```
if move->value() > best->value():  
best = move
```

TODO 5: update episode

Here is the pseudo code how I implement the update_episode part:

```
next_state_value = 0;↵
cur_value = 0;↵
↵
for (from terminal_state back to initial_state) {↵
    state& move = last_state↵
    cur_value = estimate(move.before_state());↵
    error = reward + next_state_value - cur_value;↵
    next_state_value = update(move.before_state(), a*error);↵
}↵
```

The most important thing in update_episode i think is to calculate the error.

According to the given formula, the update formula is:

$$V(S) = V(S) + a(r + V(S') - V(S))$$

So my thought is to iterate through all the moves that have been done and update the move's before_state value thus to set it as the next_state_value for next iteration, which will be:

if there are total 4 moves:

$$[s1, s1'] \rightarrow [s2, s2'] \rightarrow [s3, s3'] \rightarrow [s4, s4']$$

the update will be:

$$V(s3) \leftarrow V(s3) + a(s3.reward + V(s4) - V(s3))$$

$$V(s2) \leftarrow V(s2) + a(s2.reward + V(s3) - V(s2))$$

$$V(s1) \leftarrow V(s1) + a(s1.reward + V(s2) - V(s1))$$

The $V(s3)$ in updating $V(s2)$ is derived from the updated $V(s3)$, and $V(s2)$, $V(s1)$ likewise.

Since the terminal board's value is 0, next_state_value is initialized to 0 at first. Then iterate from state before terminal state till initial state, do update. To get the value, I simply use estimate: estimate the current move's before_state. Then calculate the error according to the given formula and thus set the updated value as the next_state_value, so that in the next iteration it can be used.